

Container Networking

Petr Horáček

What is this about?

What is this about?

- Cluster networking **evolution**

What is this about?

- Cluster networking **evolution**
- Container networking **guts**

What is this about?

- Cluster networking **evolution**
- Container networking **guts**
- **Wonderful** world of Kubernetes **networking**

What is this about?

- Cluster networking **evolution**
- Container networking **guts**
- **Wonderful** world of Kubernetes **networking**
- **Using** Kubernetes networking

Why should I care?

- Overview why is it different, how it works and how to use it

Networking Evolution

Networking Evolution

What changed since 80s?

Evolution

physical ages

Machine


Machine

Machine

Machine


Evolution

physical ages




Evolution

physical ages



Evolution

virtual ages



Evolution

cloud ages

Service

Service

Service

Service

Service

Service

Service

Service

Evolution

cloud ages



Evolution

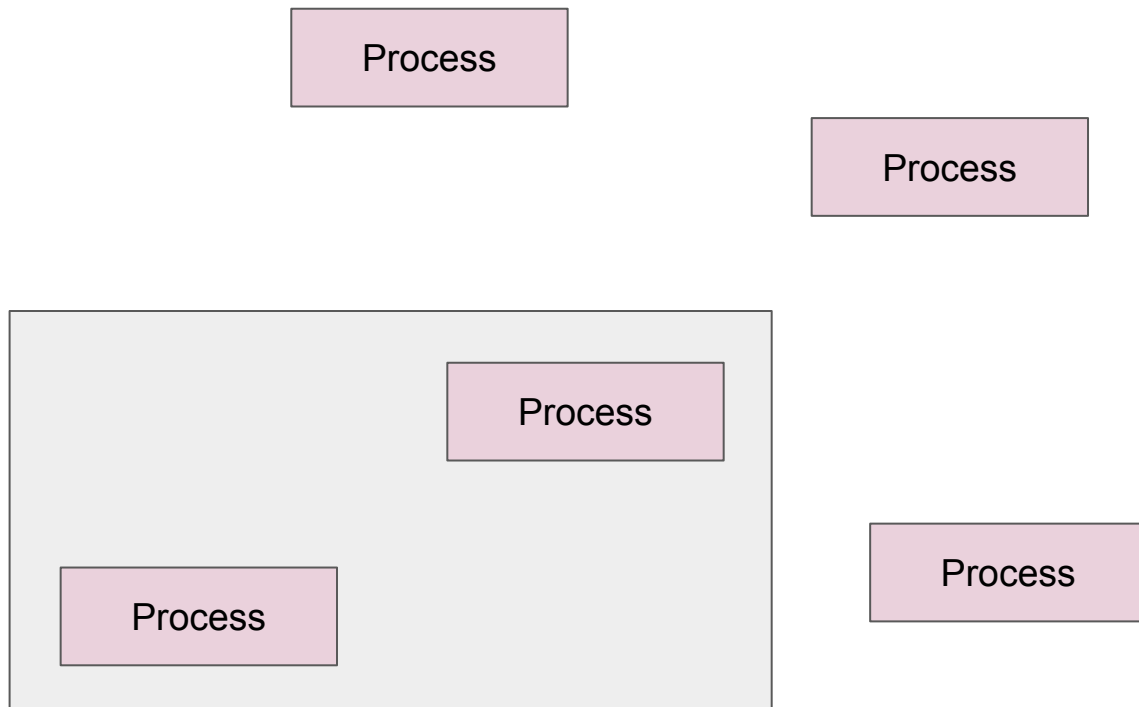
- So that's why we need different networking concept for containers

Connecting Containers

Connecting Containers

Little bit of (necessary) internals, what are containers and how do we connect them?


Container?!



Container?!


Namespaces

- resources
- filesystem
- processes
- network




Connecting

to host



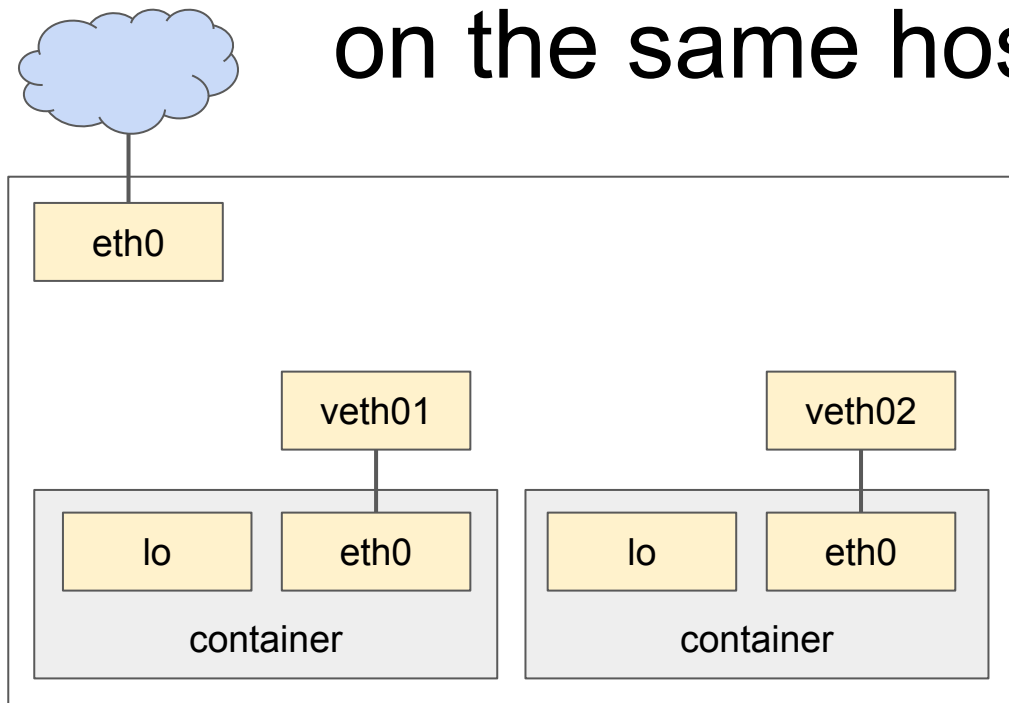
Connecting

to host



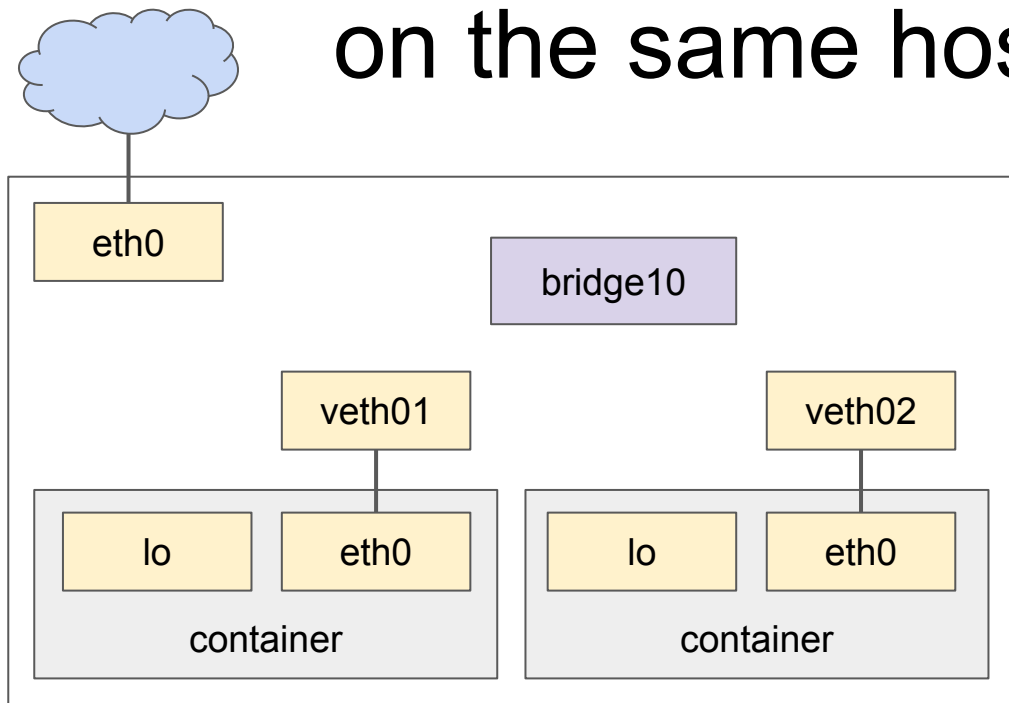
Connecting

between containers on the same host



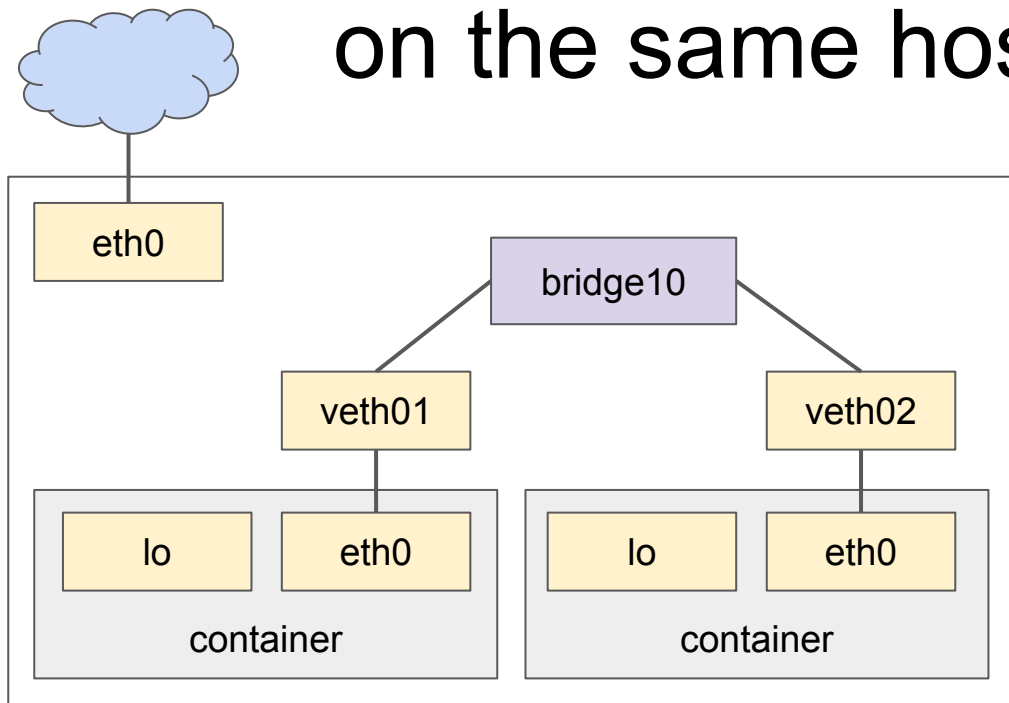
Connecting

between containers on the same host




Connecting

between containers on the same host



Connecting


to internet



Connecting

to internet

- Lies...



Connecting containers

- Using virtual interfaces to connect namespaces


Connecting Containers

Connecting Containers In Kubernetes


Connecting Containers In Kubernetes

There is a lot of networking going on in Kubernetes! It is exciting and charming in its simplicity... But it may get a little overwhelming at first.


Actually, connecting Pods, not containers




Actually, connecting Pods, not containers



Actually, connecting Pods, not containers

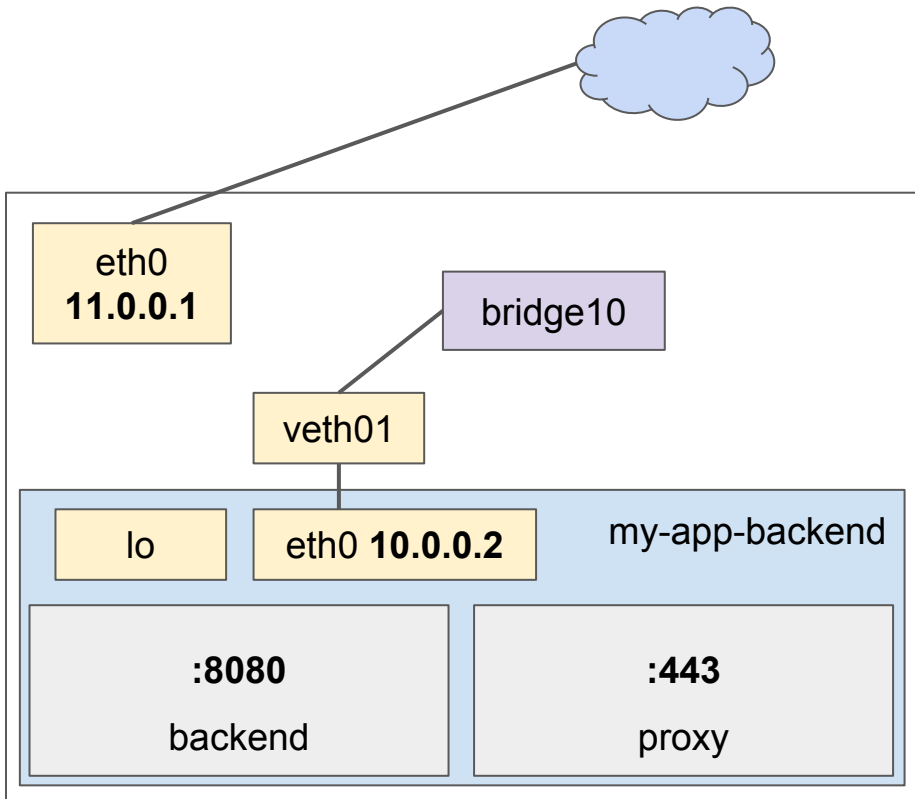


Actually, connecting Pods, not containers



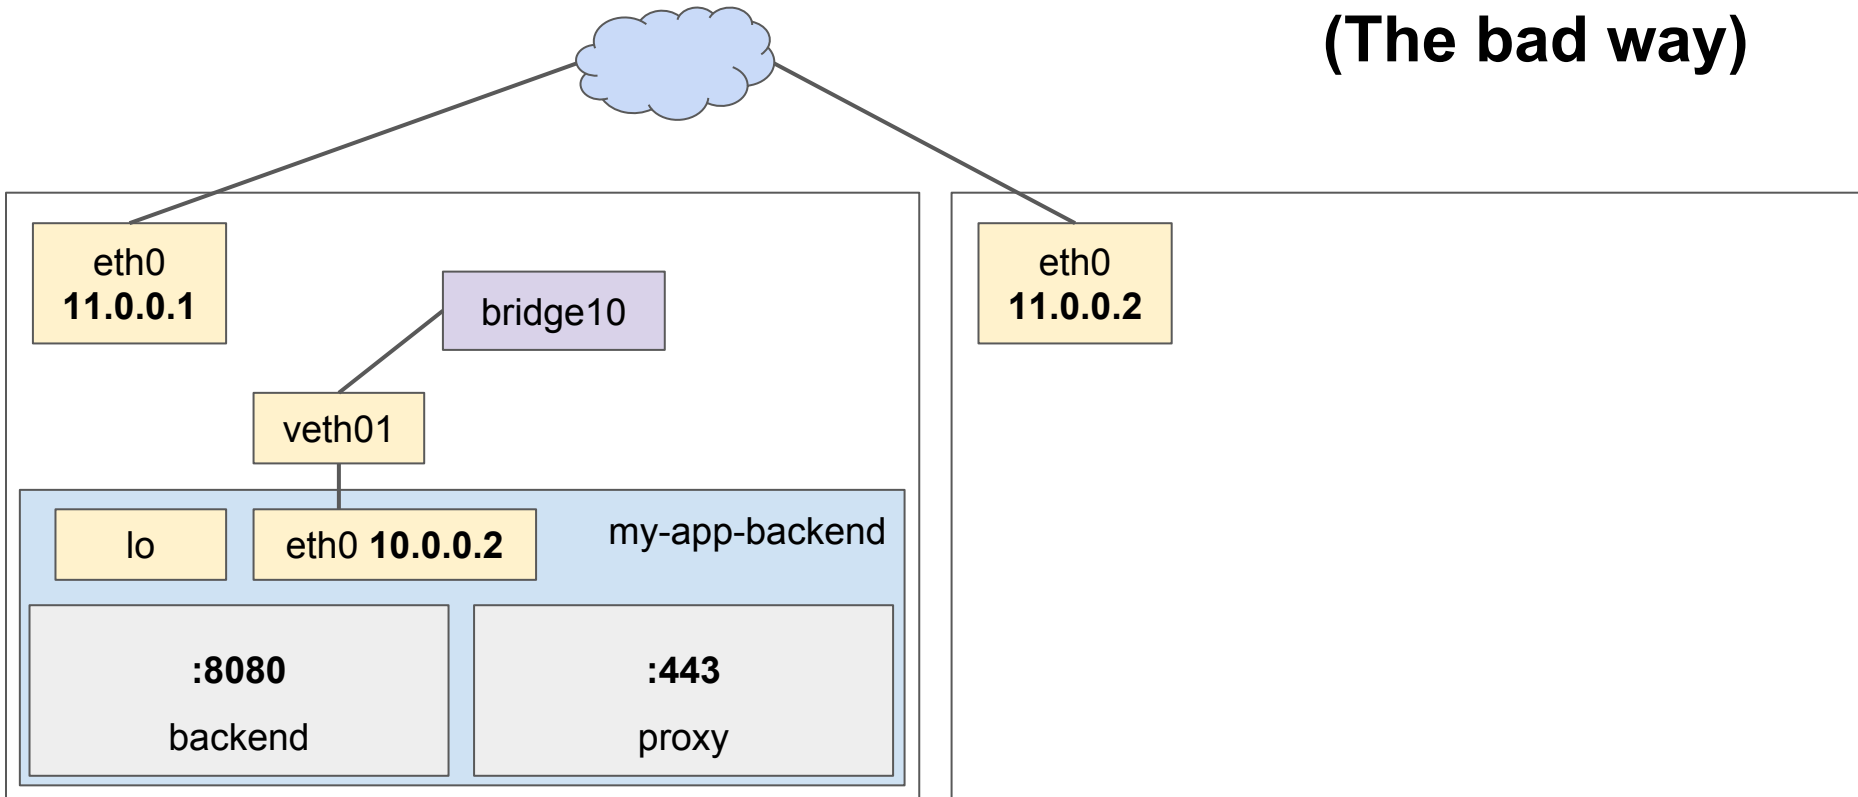
How do I access Pods on other Nodes?

(The bad way)




How do I access Pods on other Nodes?

(The bad way)



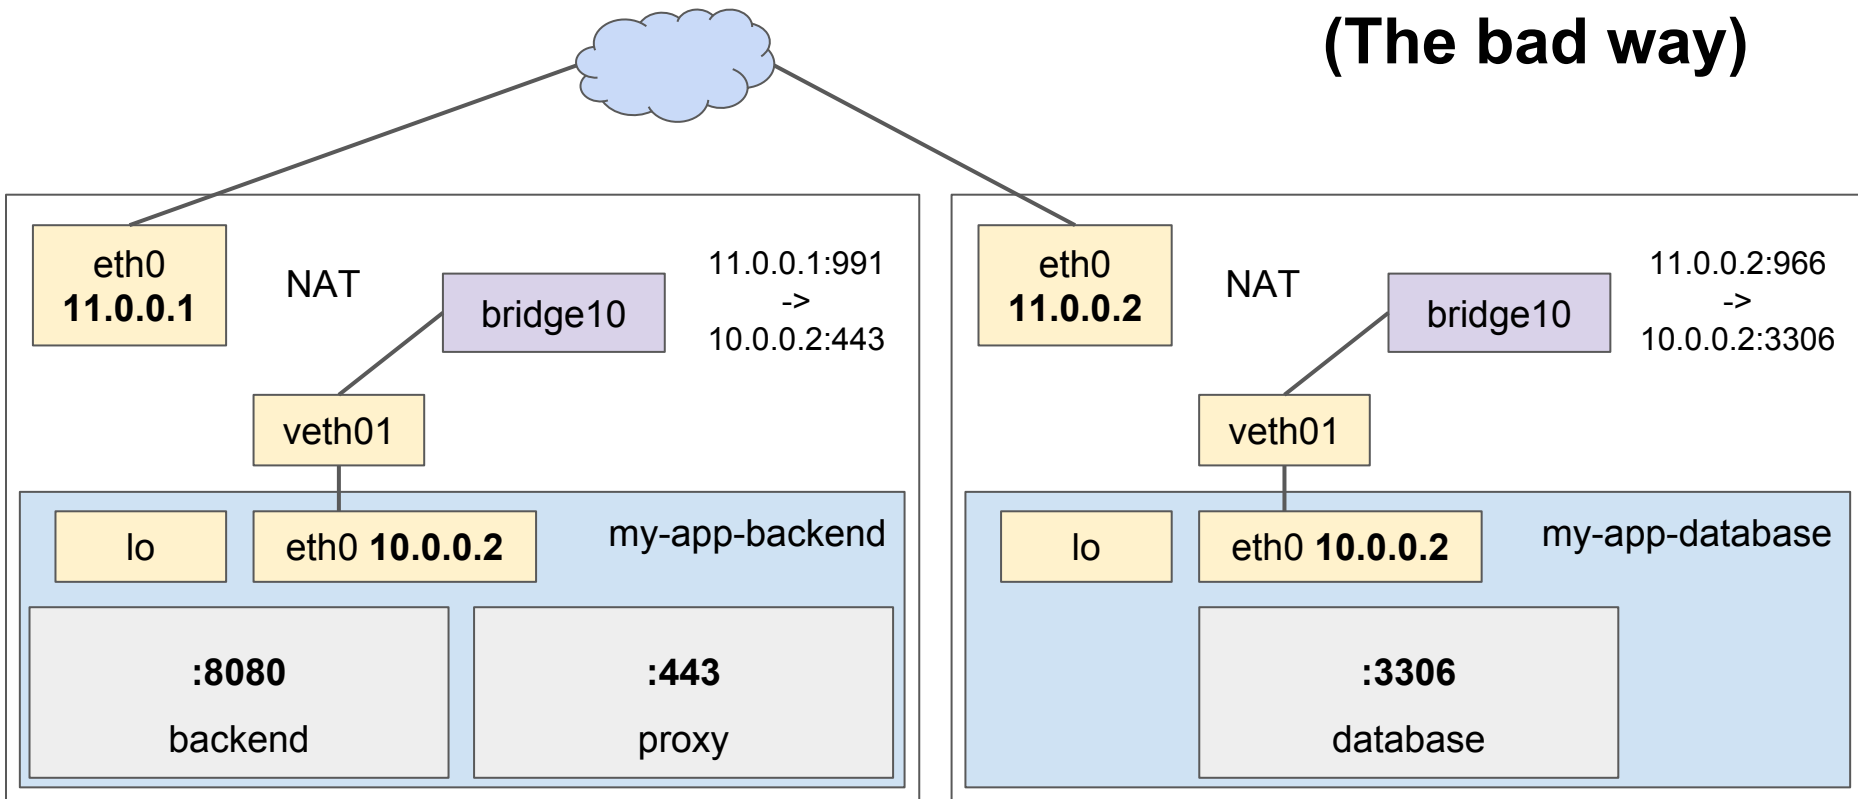
How do I access Pods on other Nodes?

(The bad way)




How do I access Pods on other Nodes?


(The bad way)



Pods can reach one another



One IP per Pod, always



Kubernetes doesn't handle networks

Kubernetes doesn't handle networks



Connecting containers in Kubernetes

- How life of a packet in Kubernetes looks like

Hands on

Hands on

Pods, Services, Endpoints, Ingress,
NetworkPolicies, ...

Let's use them (and take a closer look).

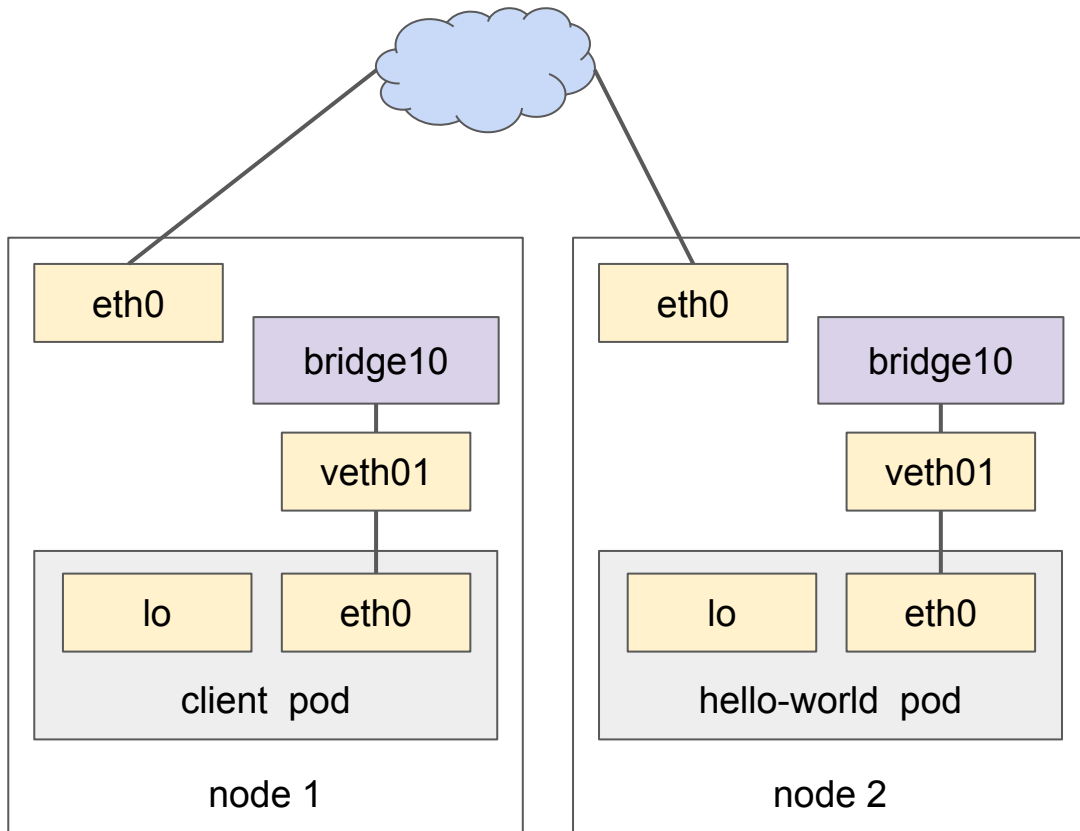
Pod again: Spec

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-world
spec:
  containers:
  - name: hello-world
    image: hello-world:1.0
    ports:
    - containerPort: 80
```

Pod again: State

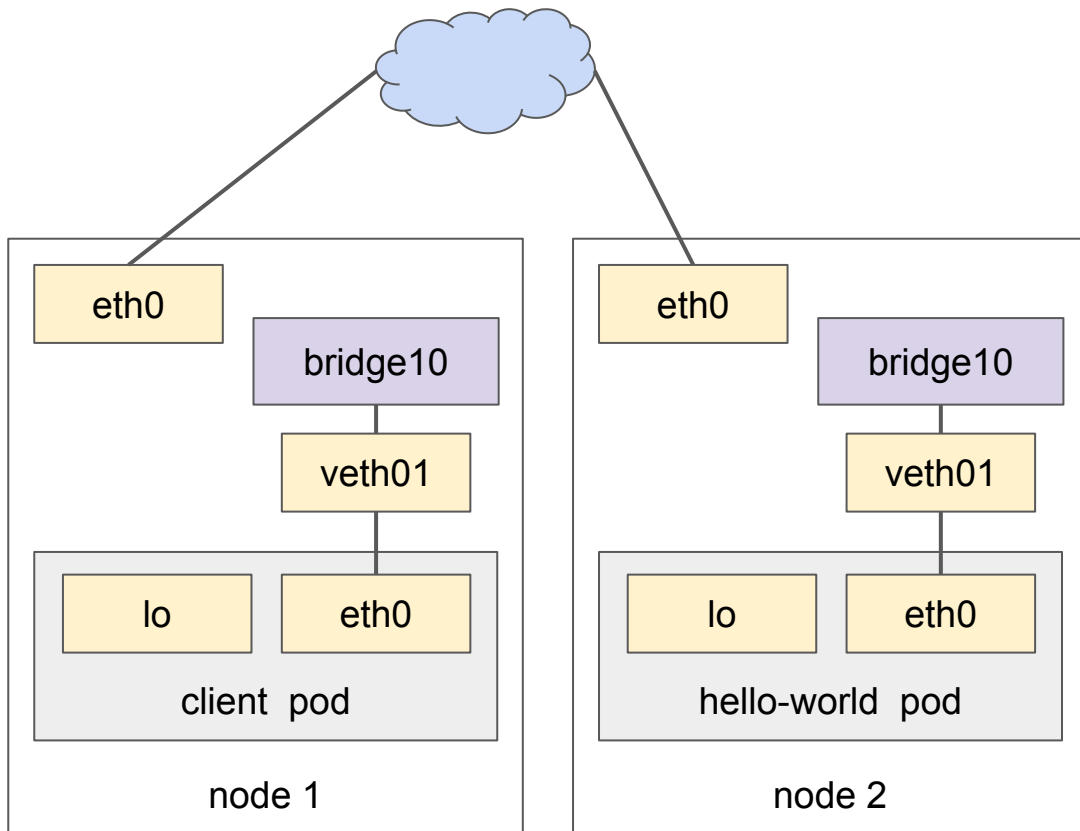
```
apiVersion: v1
kind: Pod
metadata:
  ...
  name: hello-world
  ...
spec:
  ...
status:
  ...
  podIP: 10.244.0.7
  ...
```


Pod again: Connecting



Pod again: Connecting

```
client $ curl 10.244.0.7  
"Hello World"
```




Exposing a Service: Spec

```
apiVersion: v1
kind: Service
metadata:
  name: hello-world-service
spec:
  type: ClusterIP
  ports:
    - port: 8080
      targetPort: 80
      protocol: TCP
  selector:
    app: hello-world
```

Exposing a Service: State

```
apiVersion: v1
kind: Service
metadata:
  ...
  name: hello-world-service
  ...
spec:
  clusterIP: 10.103.8.175
  ...
status:
  ...
```


Exposing a Service: Connecting



Exposing a Service: Connecting

```
client $ curl 10.103.8.175  
"Hello World"
```

```
client $ curl  
hello-world-service  
"Hello World"
```

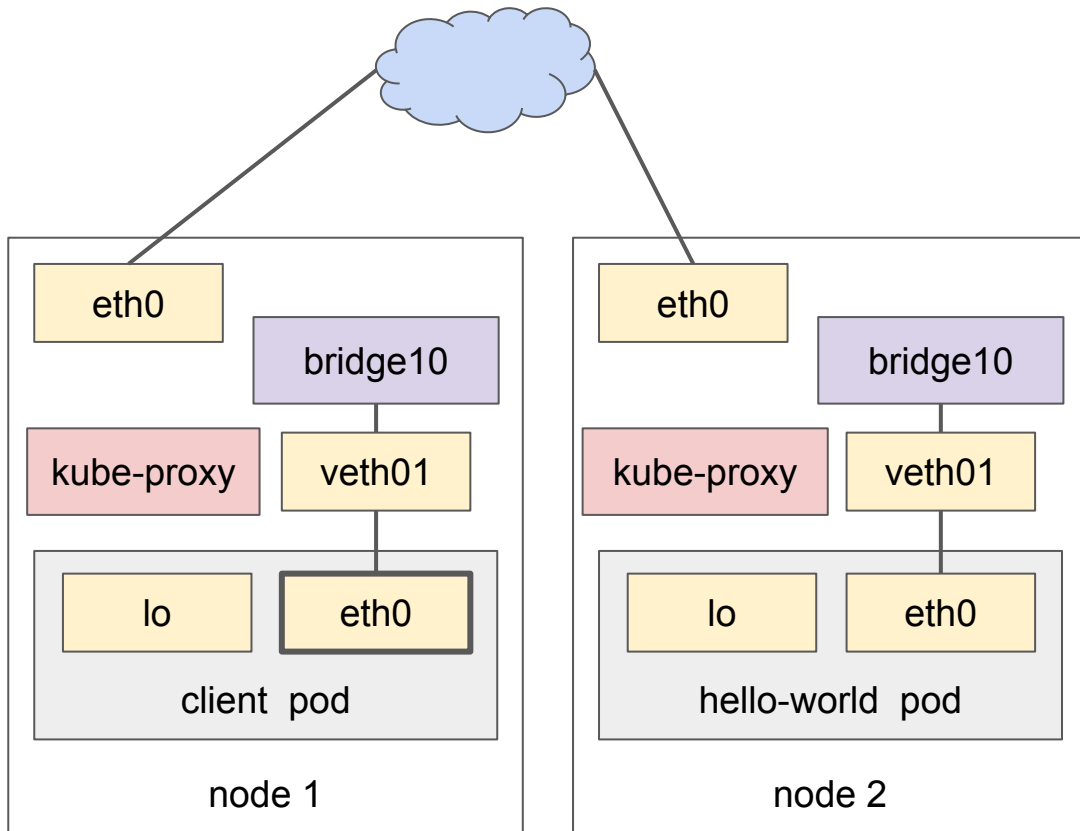


Exposing a Service: Connecting

```
client $ curl 10.103.8.175  
"Hello World"
```

```
client $ curl  
hello-world-service  
"Hello World"
```

```
From: <eth0 IP>  
To: <hello-world service IP>  
Data: ...
```

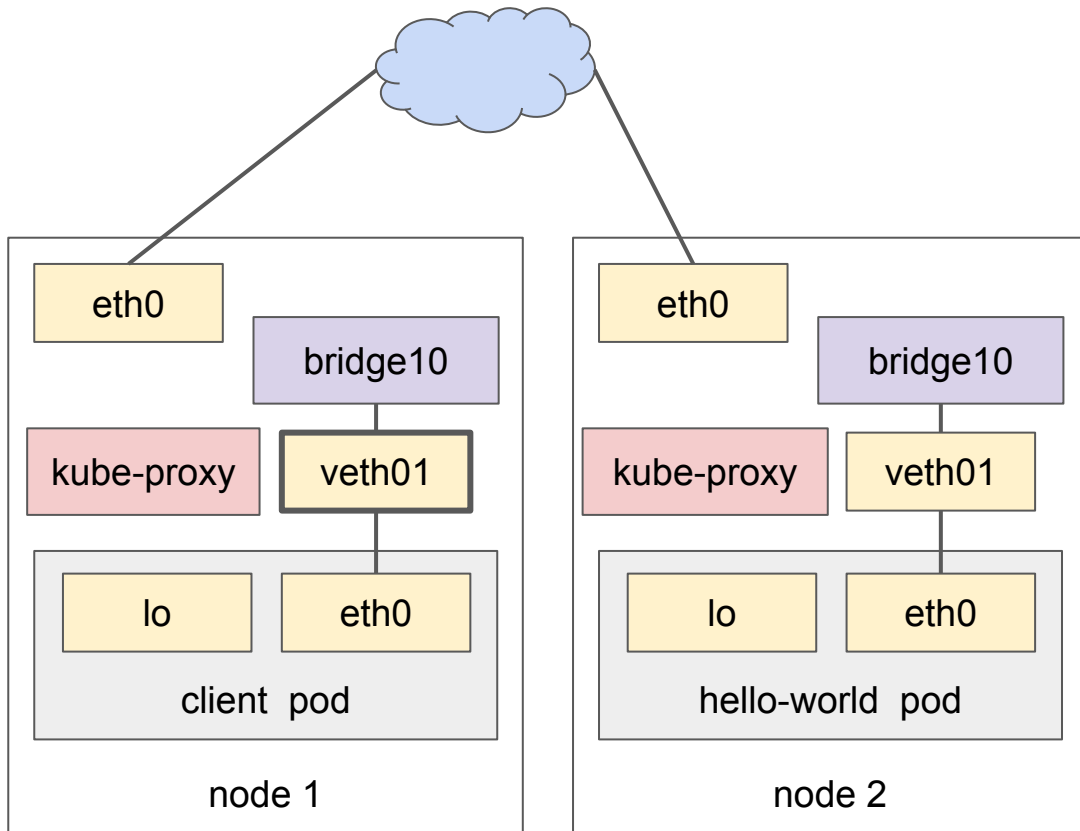


Exposing a Service: Connecting

```
client $ curl 10.103.8.175  
"Hello World"
```

```
client $ curl  
hello-world-service  
"Hello World"
```

From: <eth0 IP>
To: <hello-world service IP>
Data: ...

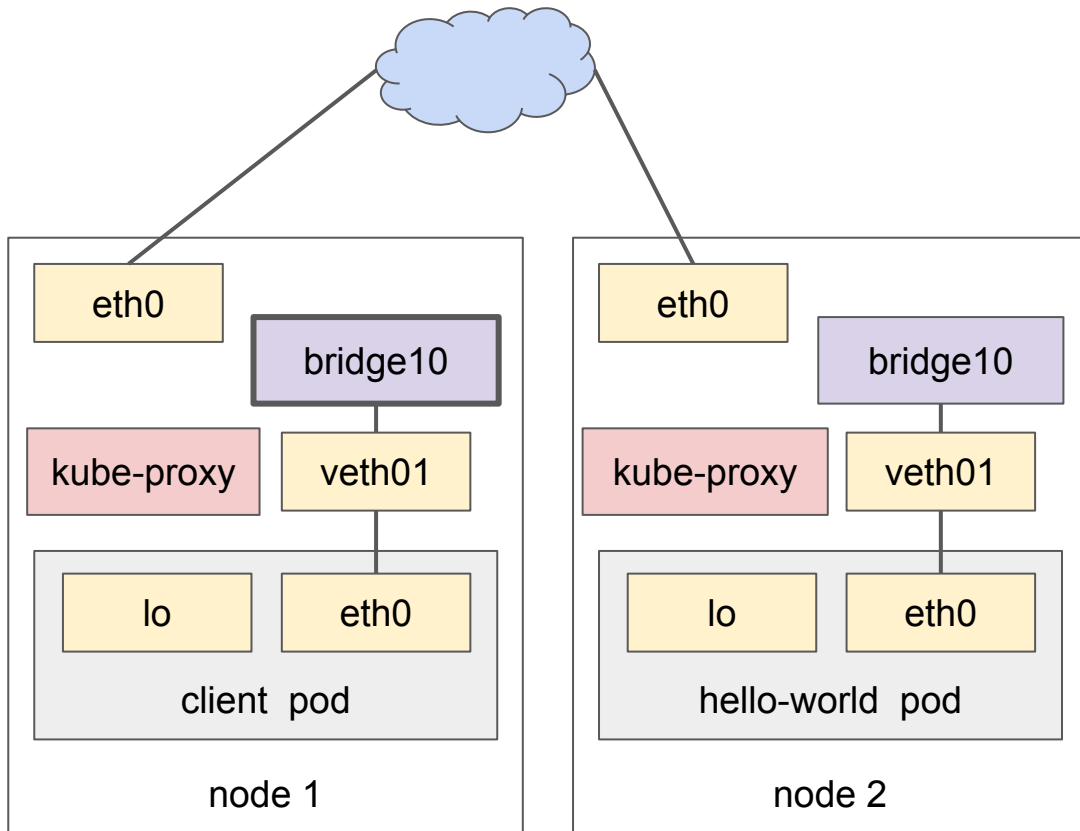


Exposing a Service: Connecting

```
client $ curl 10.103.8.175  
"Hello World"
```

```
client $ curl  
hello-world-service  
"Hello World"
```

From: <eth0 IP>
To: <hello-world service IP>
Data: ...

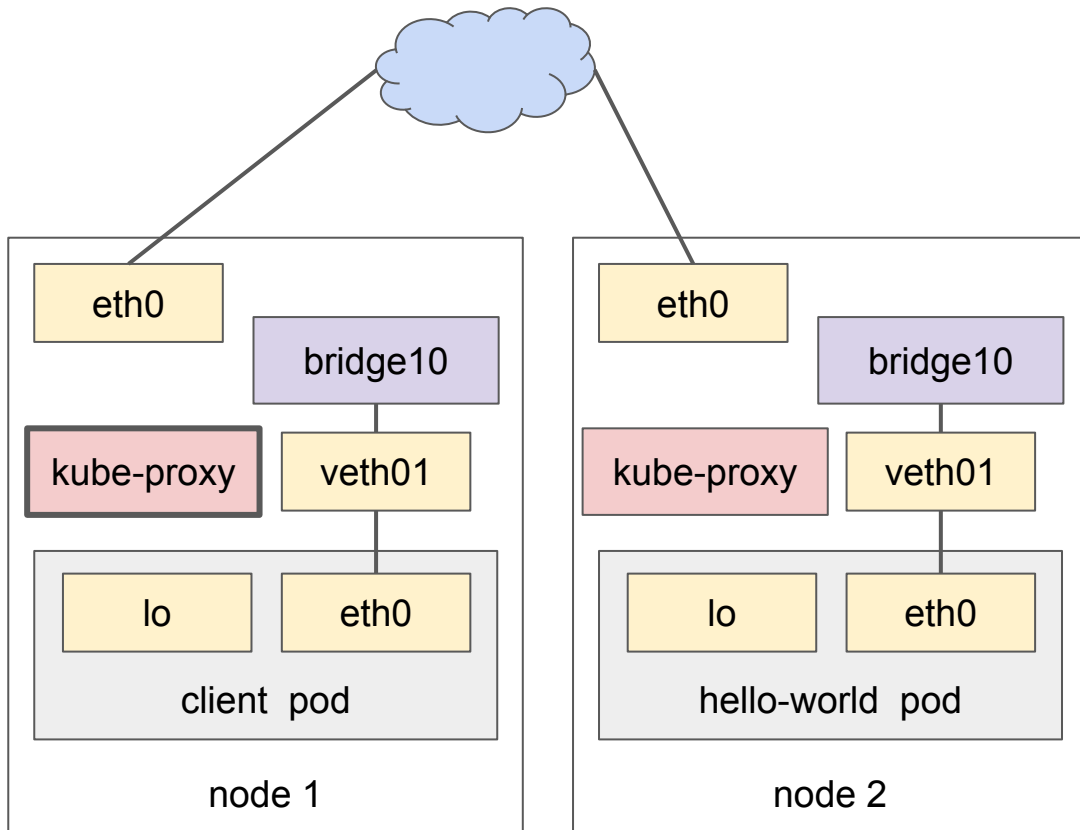


Exposing a Service: Connecting

```
client $ curl 10.103.8.175  
"Hello World"
```

```
client $ curl  
hello-world-service  
"Hello World"
```

From: <eth0 IP>
To: <hello-world service IP>
Data: ...

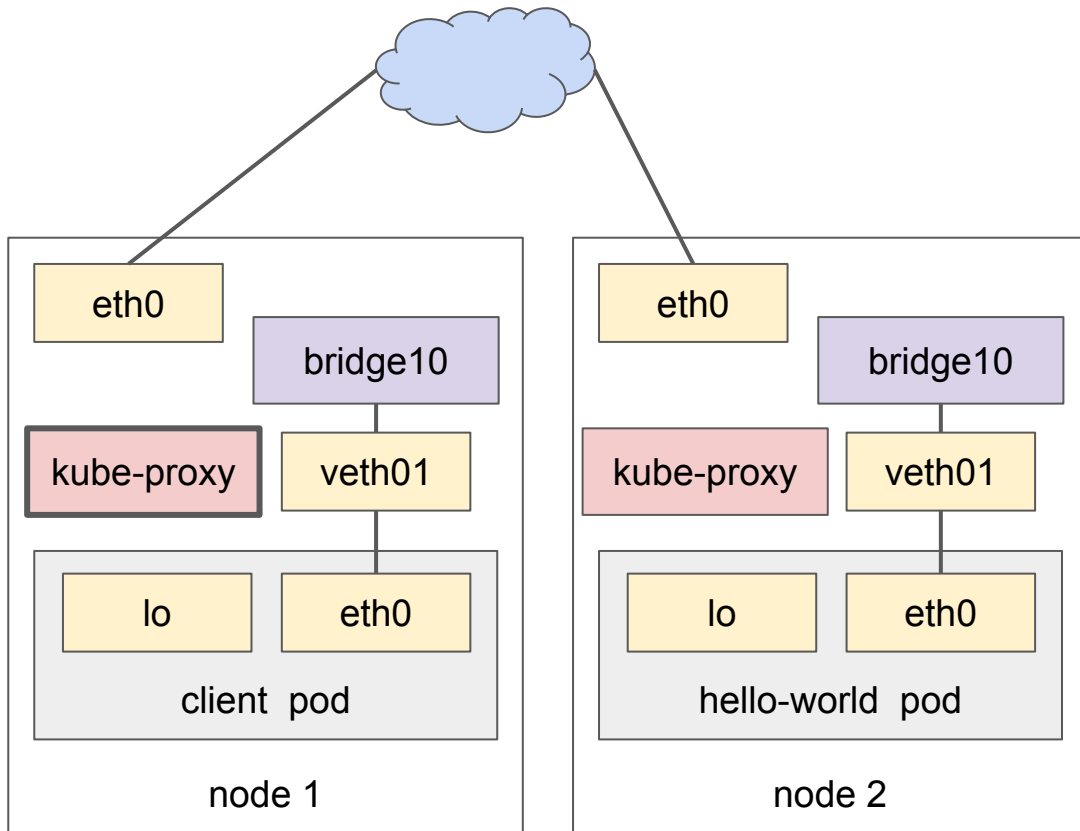


Exposing a Service: Connecting

```
client $ curl 10.103.8.175  
"Hello World"
```

```
client $ curl  
hello-world-service  
"Hello World"
```

From: <eth0 IP>
To: <random hello-world pod
IP>
Data: ...



Getting outside

- iptables NAT on each host

Getting inside

- Network load balancing (L4)
- HTTP load balancing (L7)

Getting inside to L4

Service: Spec

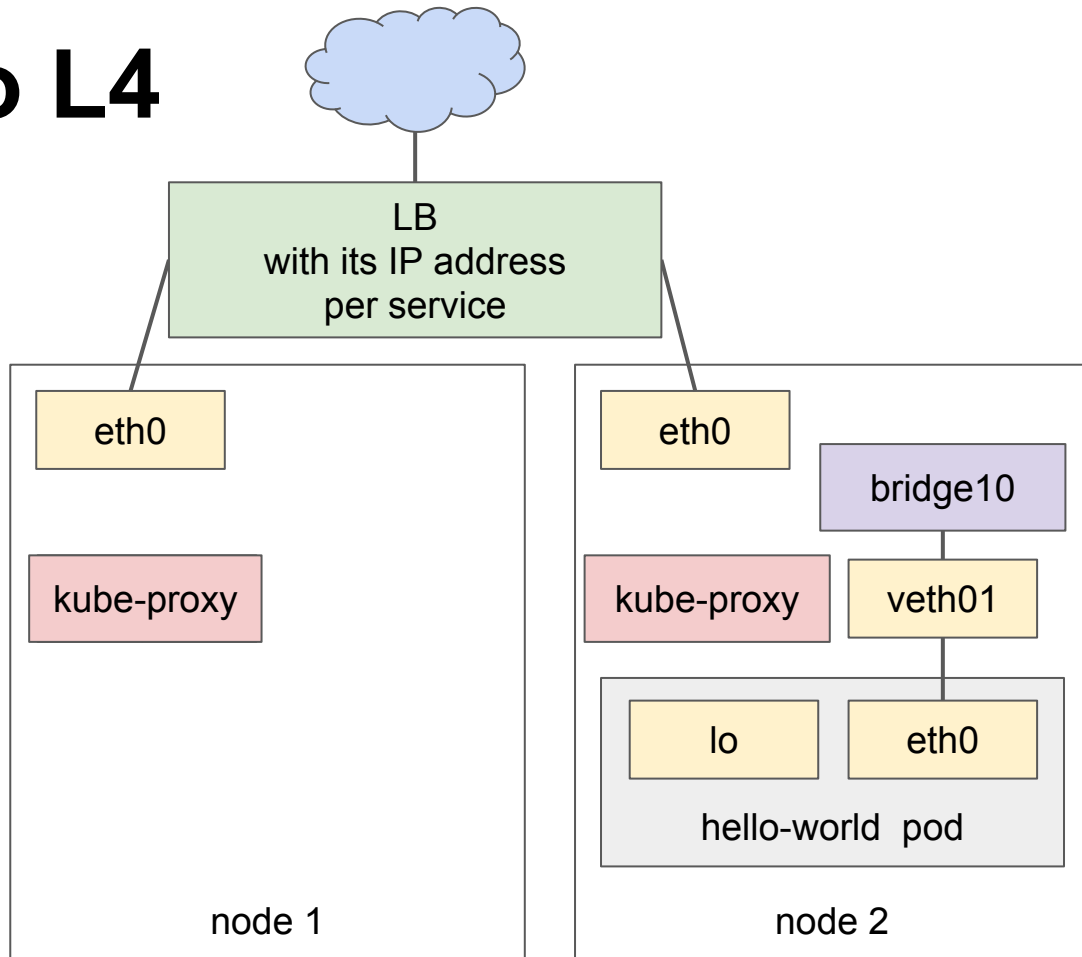
```
apiVersion: v1
kind: Service
metadata:
  name: hello-world
spec:
  type: LoadBalancer
  ports:
    - name: http
      port: 80
  selector:
    app: hello-world
```

Getting inside to L4

Service: State

```
apiVersion: v1
kind: Service
metadata:
  name: hello-world
spec:
  clusterIP: 10.103.8.176
  ...
state:
  loadBalancer:
    ingress:
      - ip: 86.105.20.11
```

Getting inside to L4

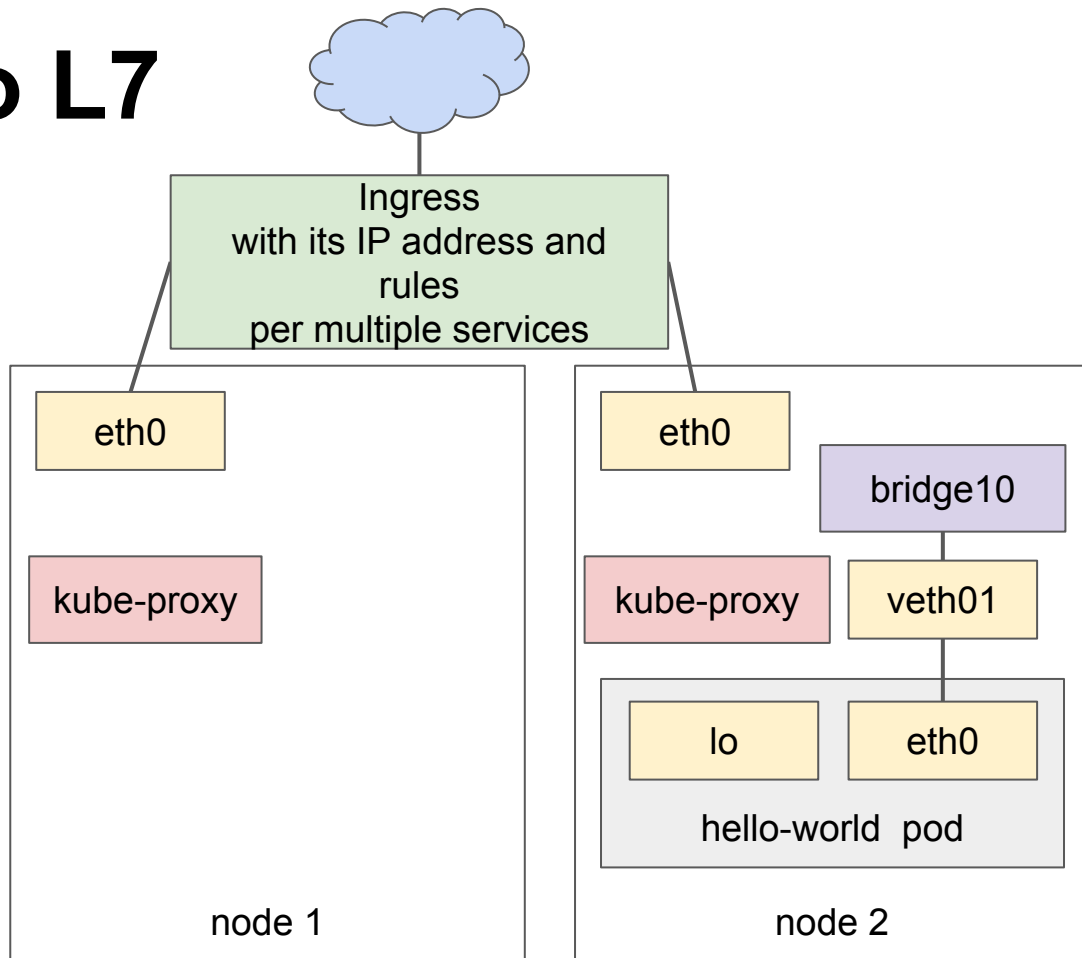


Getting inside to L7


```
spec:
  rules:
    - host: world.hello.org
      http:
        paths:
          - backend:
              serviceName: hello-world
              servicePort: 8080
    - host: darkness.hello.org
      http:
        paths:
          - path: /void/*
            backend:
              serviceName: hello-darkness
              servicePort: 8080

apiVersion:
extensions/v1beta1
kind: Ingress
metadata:
  name: hello-ingress
```


Getting inside to L7



Limiting Access



Limiting Access



Limiting Access: Spec

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
spec:
  podSelector:
    matchLabels:
      role: backend
  policyTypes:
    - Ingress
    - Egress
```

```
ingress:
- from:
  - podSelector:
      matchLabels:
        role: frontend
  ports:
    - protocol: TCP
      port: 6379
egress:
- to:
  - podSelector:
      matchLabels:
        role: database
  ports:
    - protocol: TCP
      port: 5978
```

Hands on

- Available components and how can we use them

Thanks for listening!
Bye

Check out <https://skillsmatter.com/skillscasts/10466-deep-dive-on-kubernetes-networking>